

Architecture of PFS

- Page file is a sequence of pages
 - A *page* is a basic unit of processing
 - All pages have the same length
 - A *page identifier* (abbr. *pgid*) points to the beginning of the given page (offset of a page in page file)
 - The page FS does not know anything about the contents of pages
- Architecture of Page file server
 - Page server is an Erlang process
 - * Page file server accepts requests from possibly many (ISAM) processes
 - * Process defines callback routines that define the messages (protocol of process)
 - * Process is connected by using data streams implemented in `query_node.erl` (maybe it should be renamed to `streams.erl`)
 - Binary file storage of pages
 - * N-th page is accessed by reading|writing from|to the position $n \times \text{page-size}$ in db file
 - * Pages are stored in binaries (unused fragments at the end of block)
 - * The question is weather the stream data pages are of the same size as file pages
 - * Read operation reads N pages from the given starting position in file
 - * Write operation writes N pages from the given starting position
 - * Append operation appends N pages to the end of data file
 - * Data is needed for read and write operation is transfered via data streams
 - * Protocol thus require the completion message for read and write (in opposite directions)
 - PFS is linked to a client vie I/O data streams
 - * Input/output data of write/read operations is obtained via data streams
 - * Data streams are composed of data messages that contain up to `TRIPLES_IN_PAGE` triples (or less)
 - * Reading/writing data messages/triples from/to a stream
 - Processing unit is either a data message or a triple
 - Stream type is defined on initialization of a named queue

- Requests are placed in a queue and served one by one
 - * Pid of the client process is stored for each request
 - * Request to read N pages is completed after all the pages are read and sent to client
 - * Request to write N pages starts after complete data has been transfered
 - Data can be stored in a map that maps Pids to lists of collected data pages
 - * Each request can process (read or write) a chunk of data
 - A chunk of data is defined by the number of pages
 - After a chunk is processed the state is stored in request and it is put back at the end of queue
 - This implements a kind of round-robin algorithm
 - All other request do not freeze if a large request is being processed
 - * (to-do) Does it make sense to have sessions (with a given process pid)?
- (expand) A cache is part of PFS
 - * Page are read into buffer pool
 - * LRU page replacement strategy is used
- Page file server interface
 - { data_read, pgid, N }
 - * Reads a sequence of N pages starting at the page pgid
 - * Read data pages are sent to the client process via data streams
 - { data_read_end }
 - * Signals the completion of data_read operation
 - * Number of pages sent to client is N
 - { data_write, pgid, N }
 - * Writes a sequence of N pages to the db file
 - * Data pages to write are received from a client process via data streams
 - { data_write_end }
 - * Signals the completion of data_write or data_append operation
 - * Number of pages received from client has to be equal N
 - { data_append, N } }
 - * Writes a sequence of N pages to the end of data file
 - * Data pages to write are received from a client process via data streams